

1 Business Problem

Your company is expanding in to new industries to diversify its portfolio. Specifically, they are interested in purchasing and operating airplanes for commercial and private enterprises, but do not know anything about the potential risks of aircraft. You are charged with determining which aircraft are the lowest risk for the company to start this new business endeavor.

The Objective is: To determine risk evaluation of different aircrafts focusing on North and South America countries with most accidents and incidents

2 1. Data Loading and Initial Exploration

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: #load the dataset
file_path = "/home/khalid-omar/Documents/aviation/AviationData.csv"
df = pd.read_csv(file_path, encoding='latin1', low_memory=False)
print("First few rows of the dataset:")
df.head()
```

First few rows of the dataset:

Out[2]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Col
0	20001218X45444	Accident	SEA87LA080	1948-10-24	MOOSE CREEK, ID	
1	20001218X45447	Accident	LAX94LA336	1962-07-19	BRIDGEPORT, CA	
2	20061025X01555	Accident	NYC07LA005	1974-08-30	Saltville, VA	
3	20001218X45448	Accident	LAX96LA321	1977-06-19	EUREKA, CA	
4	20041105X01764	Accident	CHI79FA064	1979-08-02	Canton, OH	

5 rows × 31 columns

Checking Dataset Overview:

In [3]: `df.tail()`

Out[3]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Col
88884	20221227106491	Accident	ERA23LA093	2022-12-26	Annapolis, MD	
88885	20221227106494	Accident	ERA23LA095	2022-12-26	Hampton, NH	
88886	20221227106497	Accident	WPR23LA075	2022-12-26	Payson, AZ	
88887	20221227106498	Accident	WPR23LA076	2022-12-26	Morgan, UT	
88888	20221230106513	Accident	ERA23LA097	2022-12-29	Athens, GA	

5 rows × 31 columns



In [4]: `df.shape`

Out[4]: (88889, 31)

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 88889 entries, 0 to 88888
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Event.Id                             88889 non-null  object
1   Investigation.Type                   88889 non-null  object
2   Accident.Number                     88889 non-null  object
3   Event.Date                          88889 non-null  object
4   Location                            88837 non-null  object
5   Country                            88663 non-null  object
6   Latitude                           34382 non-null  object
7   Longitude                           34373 non-null  object
8   Airport.Code                       50249 non-null  object
9   Airport.Name                       52790 non-null  object
10  Injury.Severity                     87889 non-null  object
11  Aircraft.damage                     85695 non-null  object
12  Aircraft.Category                   32287 non-null  object
13  Registration.Number                 87572 non-null  object
14  Make                               88826 non-null  object
15  Model                              88797 non-null  object
16  Amateur.Built                      88787 non-null  object
17  Number.of.Engines                  82805 non-null  float64
18  Engine.Type                        81812 non-null  object
19  FAR.Description                    32023 non-null  object
20  Schedule                          12582 non-null  object
21  Purpose.of.flight                  82697 non-null  object
22  Air.carrier                        16648 non-null  object
23  Total.Fatal.Injuries                77488 non-null  float64
24  Total.Serious.Injuries              76379 non-null  float64
25  Total.Minor.Injuries                76956 non-null  float64
26  Total.Uninjured                    82977 non-null  float64
27  Weather.Condition                   84397 non-null  object
28  Broad.phase.of.flight               61724 non-null  object
29  Report.Status                       82508 non-null  object
30  Publication.Date                    75118 non-null  object
dtypes: float64(5), object(26)
memory usage: 21.0+ MB
```

In [6]: `df.info(verbose=False)`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 88889 entries, 0 to 88888
Columns: 31 entries, Event.Id to Publication.Date
dtypes: float64(5), object(26)
memory usage: 21.0+ MB
```

```
In [7]: # Check the columns in the DataFrame
print(df.columns)
```

```
Index(['Event.Id', 'Investigation.Type', 'Accident.Number', 'Event.Date',
      'Location', 'Country', 'Latitude', 'Longitude', 'Airport.Code',
      'Airport.Name', 'Injury.Severity', 'Aircraft.damage',
      'Aircraft.Category', 'Registration.Number', 'Make', 'Model',
      'Amateur.Built', 'Number.of.Engines', 'Engine.Type', 'FAR.Description',
      'Schedule', 'Purpose.of.flight', 'Air.carrier', 'Total.Fatal.Injuries',
      'Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Uninjured',
      'Weather.Condition', 'Broad.phase.of.flight', 'Report.Status',
      'Publication.Date'],
      dtype='object')
```

3 2. Data Cleaning and Filtering

```
In [8]: df['Country'].value_counts()[0:25]
```

```
Out[8]: United States      82248
Brazil                    374
Canada                    359
Mexico                    358
United Kingdom            344
Australia                  300
France                     236
Spain                      226
Bahamas                    216
Germany                    215
Colombia                   193
South Africa               129
Japan                      126
Venezuela                  121
Italy                      114
Argentina                  112
Indonesia                  110
India                       96
Peru                       93
Russia                     91
ATLANTIC OCEAN             81
Ireland                    77
Puerto Rico               71
Dominican Republic         68
Guatemala                  67
Name: Country, dtype: int64
```

Filtering Countries:

```
In [9]: # List of countries to keep
countries_to_keep = [
    "United States", "Brazil", "Canada", "Mexico",
    "Bahamas", "Colombia", "Venezuela", "Argentina",
    "Peru", "Puerto Rico"
]

# Filter the DataFrame to only include rows where 'Country' is in the
df = df[df['Country'].isin(countries_to_keep)]

# Optionally, sort the DataFrame by 'Country'
df = df.sort_values(by='Country')

# Check the result
print(df['Country'].value_counts())
```

```
United States    82248
Brazil           374
Canada           359
Mexico           358
Bahamas          216
Colombia         193
Venezuela        121
Argentina        112
Peru              93
Puerto Rico      71
Name: Country, dtype: int64
```

In [10]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 84145 entries, 57740 to 71219
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Event.Id                             84145 non-null  object
1   Investigation.Type                    84145 non-null  object
2   Accident.Number                      84145 non-null  object
3   Event.Date                           84145 non-null  object
4   Location                             84132 non-null  object
5   Country                             84145 non-null  object
6   Latitude                             33032 non-null  object
7   Longitude                            33022 non-null  object
8   Airport.Code                         49526 non-null  object
9   Airport.Name                         51996 non-null  object
10  Injury.Severity                      83880 non-null  object
11  Aircraft.damage                      81900 non-null  object
12  Aircraft.Category                    29319 non-null  object
13  Registration.Number                  83785 non-null  object
14  Make                                 84115 non-null  object
15  Model                                84096 non-null  object
16  Amateur.Built                       84111 non-null  object
17  Number.ofEngines                     81139 non-null  float64
18  Engine.Type                          80019 non-null  object
19  FAR.Description                      29356 non-null  object
20  Schedule                             10865 non-null  object
21  Purpose.of.flight                    80782 non-null  object
22  Air.carrier                          15164 non-null  object
23  Total.Fatal.Injuries                 73321 non-null  float64
24  Total.Serious.Injuries               72467 non-null  float64
25  Total.Minor.Injuries                 73102 non-null  float64
26  Total.Uninjured                      78897 non-null  float64
27  Weather.Condition                    82541 non-null  object
28  Broad.phase.of.flight                61259 non-null  object
29  Report.Status                        80449 non-null  object
30  Publication.Date                     71165 non-null  object
dtypes: float64(5), object(26)
memory usage: 20.5+ MB
```

Handling Missing Data:

```
In [11]: # Check for missing values
missing_values = df.isnull().sum()
missing_values
```

```
Out[11]: Event.Id          0
Investigation.Type        0
Accident.Number           0
Event.Date                0
Location                 13
Country                   0
Latitude                 51113
Longitude                 51123
Airport.Code              34619
Airport.Name              32149
Injury.Severity           265
Aircraft.damage           2245
Aircraft.Category         54826
Registration.Number        360
Make                      30
Model                     49
Amateur.Built             34
Number.of.Engines         3006
Engine.Type               4126
FAR.Description           54789
Schedule                  73280
Purpose.of.flight         3363
Air.carrier               68981
Total.Fatal.Injuries      10824
Total.Serious.Injuries    11678
Total.Minor.Injuries      11043
Total.Uninjured           5248
Weather.Condition         1604
Broad.phase.of.flight     22886
Report.Status             3696
Publication.Date          12980
dtype: int64
```

Dropping Unnecessary Columns:

```
In [12]: # List of columns to drop
columns_to_drop = ['Longitude', 'Latitude', 'Airport.Name', 'Aircraft.

# Drop the specified columns from the DataFrame
df.drop(columns=columns_to_drop, inplace=True)

# Check the updated DataFrame to confirm the changes
print(df.head())
```


	Event.Id	Investigation.Type	Accident.Number	Event.Date
\				
57740	20041109X01789	Accident	MIA05WA026	2004-10-28
36546	20001206X02782	Accident	MIA95WA047	1994-12-21
48270	20001212X20889	Accident	MIA00WA142	2000-04-29
66116	20090414X52749	Incident	DCA09WA044	2009-03-25
49930	20010130X00371	Accident	DCA01WA013	2001-01-09

	Location	Country	Airport.Code	Injury.Severit
y \				
57740	SAN FERNANDO, ARGENTINA	Argentina	NaN	Non-Fata
l				
36546	BUENOS AIRES, ARGENTINA	Argentina	NaN	Non-Fata
l				
48270	CHASCOMUS, ARGENTINA	Argentina	NaN	Non-Fata
l				
66116	Trelew, Argentina	Argentina	NaN	Non-Fata
l				
49930	BUENOS AIRES, ARGENTINA	Argentina	NaN	Non-Fata
l				

	Aircraft.damage	Registration.Number	...	Purpose.of.flight	Air.
carrier \					
57740	Substantial	NaN	...	Personal	
NaN					
36546	Substantial	N747E	...	Unknown	
NaN					
48270	Substantial	N156P	...	Business	
NaN					
66116	NaN	LV-VBZ	...	NaN	
NaN					
49930	Substantial	NaN	...	NaN	
NaN					

	Total.Fatal.Injuries	Total.Serious.Injuries	Total.Minor.Injuri
es \			
57740	NaN	NaN	
3.0			
36546	0.0	0.0	
0.0			
48270	0.0	0.0	
4.0			
66116	0.0	0.0	
0.0			
49930	NaN	NaN	N
aN			

	Total.Uninjured	Weather.Condition	Broad.phase.of.flight	Repor
t.Status \				
57740	NaN	VMC	NaN	
Foreign				
36546	2.0	VMC	NaN	
Foreign				
48270	0.0	VMC	NaN	
Foreign				
66116	71.0	NaN	NaN	
.				

49930	144.0	NaN	NaN
Foreign			

	Publication.Date
57740	09-11-2004
36546	30-12-1994
48270	15-05-2000
66116	03-11-2020
49930	NaN

[5 rows x 25 columns]

```
In [13]: # Check for missing values
missing_values = df.isnull().sum()
missing_values
```

```
Out[13]: Event.Id          0
Investigation.Type        0
Accident.Number           0
Event.Date                0
Location                  13
Country                   0
Airport.Code              34619
Injury.Severity           265
Aircraft.damage           2245
Registration.Number       360
Make                       30
Model                     49
Amateur.Built             34
Number.of.Engines         3006
Engine.Type               4126
Purpose.of.flight         3363
Air.carrier               68981
Total.Fatal.Injuries      10824
Total.Serious.Injuries    11678
Total.Minor.Injuries      11043
Total.Uninjured           5248
Weather.Condition         1604
Broad.phase.of.flight     22886
Report.Status             3696
Publication.Date          12980
dtype: int64
```

4 3. Imputation and Outlier Handling

Imputation of Missing Values:

```
In [14]: # Impute missing values in categorical columns with the mode (most frequent)
categorical_columns = df.select_dtypes(include=['object']).columns

for col in categorical_columns:
    df[col].fillna(df[col].mode()[0], inplace=True)
```

```
In [15]: # Check for missing values after imputation
missing_values_after_imputation = df.isnull().sum()
print(missing_values_after_imputation)
```

```
Event.Id                0
Investigation.Type      0
Accident.Number         0
Event.Date              0
Location                0
Country                 0
Airport.Code            0
Injury.Severity         0
Aircraft.damage         0
Registration.Number     0
Make                    0
Model                   0
Amateur.Built           0
Number.ofEngines        3006
Engine.Type             0
Purpose.of.flight       0
Air.carrier             0
Total.Fatal.Injuries    10824
Total.Serious.Injuries  11678
Total.Minor.Injuries    11043
Total.Uninjured         5248
Weather.Condition       0
Broad.phase.of.flight   0
Report.Status           0
Publication.Date        0
dtype: int64
```

```
In [16]: # Descriptive statistics
print(df.describe())
```

	Number.of.Engines	Total.Fatal.Injuries	Total.Serious.Injuries
count	81139.000000	73321.000000	72467.000000
mean	1.138035	0.478240	0.26236
std	0.431596	3.205397	1.20144
min	0.000000	0.000000	0.000000
25%	1.000000	0.000000	0.000000
50%	1.000000	0.000000	0.000000
75%	1.000000	0.000000	0.000000
max	8.000000	265.000000	137.000000

	Total.Minor.Injuries	Total.Uninjured
count	73102.000000	78897.000000
mean	0.341413	4.492858
std	1.936168	24.482405
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	1.000000
75%	0.000000	2.000000
max	380.000000	699.000000

```
In [17]: print(df.columns)
```

```
Index(['Event.Id', 'Investigation.Type', 'Accident.Number', 'Event.Date',
      'Location', 'Country', 'Airport.Code', 'Injury.Severity',
      'Aircraft.damage', 'Registration.Number', 'Make', 'Model',
      'Amateur.Built', 'Number.of.Engines', 'Engine.Type',
      'Purpose.of.flight', 'Air.carrier', 'Total.Fatal.Injuries',
      'Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Uninjured',
      'Weather.Condition', 'Broad.phase.of.flight', 'Report.Status',
      'Publication.Date'],
      dtype='object')
```

Filtering the Dataset to ensure the dataset includes only:

Relevant countries (already filtered earlier). Accidents and incidents in the Investigation.Type column. Aircraft types with sufficient representation (remove rare types to avoid statistical noise).

```
In [18]: # Keep only rows where 'Investigation.Type' is 'Accident' or 'Incident'
df = df[df['Investigation.Type'].isin(['Accident', 'Incident'])]

# Check the most common aircraft types and keep the top 10
top_aircrafts = df['Aircraft.damage'].value_counts().head(10).index
df = df[df['Aircraft.damage'].isin(top_aircrafts)]

# View filtered dataset
print(df['Aircraft.damage'].value_counts())
```

```
Substantial    64739
Destroyed      17092
Minor          2247
Unknown         67
Name: Aircraft.damage, dtype: int64
```

5 4. Risk Metrics

Fatality and Survivability Rates:

Evaluating the Risk Metrics Defining risk as the relationship between accidents/incidents and fatalities/uninjuries. Key metrics to calculate:

Fatality Rate: Average number of fatalities per aircraft category. Survivability Rate: Proportion of uninjured passengers.

```
In [19]: # Calculate fatality rate by Aircraft.damage
fatality_rate = df.groupby('Aircraft.damage')['Total.Fatal.Injuries'].

# Calculate survivability rate by Aircraft.damage
df['Survivability.Rate'] = df['Total.Uninjured'] / (df['Total.Fatal.In
survivability_rate = df.groupby('Aircraft.damage')['Survivability.Rate

# Combine risk metrics into a single DataFrame
risk_metrics = pd.DataFrame({
    'Fatality Rate': fatality_rate,
    'Survivability Rate': survivability_rate
})

print(risk_metrics)
```

Aircraft.damage	Fatality Rate	Survivability Rate
Destroyed	1.683506	0.124768
Minor	0.142857	0.903558
Substantial	0.127168	0.701343
Unknown	0.388060	0.571038

Outlier Removal Using Z-scores:

```
In [20]: # Example numerical columns
numerical_columns = ['Total.Fatal.Injuries', 'Total.Uninjured', 'Total

# Creating a copy of the dataframe for cleaning
df_cleaned_numpy = df.copy()

# Removing outliers based on Z-score (threshold = 3)
for col in numerical_columns:
    # Calculate the Z-scores for the column
    z_scores = np.abs((df_cleaned_numpy[col] - np.mean(df_cleaned_numpy

    # Keep only rows where the Z-score is less than 3 (no outliers)
    df_cleaned_numpy = df_cleaned_numpy[z_scores < 3]

print(f"Data shape after Z-score outlier removal: {df_cleaned_numpy.sh
```

Data shape after Z-score outlier removal: (68625, 26)

Highlights:

This provides a good basis to assess safety metrics across various aircraft damage types.

```
In [21]: # Remove duplicates
df.drop_duplicates(inplace=True)
```

```
In [22]: # Fill missing values with the mean, excluding non-numeric columns (li
df.fillna(df.mean(numeric_only=True), inplace=True)
```

Grouping and Aggregation

Average Injuries by Investigation Type:

```
In [23]: # Grouping and Aggregation
# Grouping data by 'Investigation.Type' and calculating the mean of 'T
print("\nAverage Total.Uninjured by Investigation Type:")
print(df.groupby('Investigation.Type')['Total.Uninjured'].mean())
```

```
Average Total.Uninjured by Investigation Type:
Investigation.Type
Accident      2.929347
Incident     52.687903
Name: Total.Uninjured, dtype: float64
```

```
In [24]: # Average Total.Uninjured by Investigation Type
print(df.groupby('Investigation.Type')['Total.Uninjured'].mean())
```

```
Investigation.Type
Accident      2.929347
Incident     52.687903
Name: Total.Uninjured, dtype: float64
```

```
In [25]: # Ensure column names are stripped of whitespace
df.columns = df.columns.str.strip()

# Grouping data by 'Investigation.Type' and calculating the mean of 'Total.Fatal.Injuries'
if 'Investigation.Type' in df.columns and 'Total.Fatal.Injuries' in df.columns:
    avg_fatal_injuries = df.groupby('Investigation.Type')['Total.Fatal.Injuries'].mean()
    print("\nAverage Total.Fatal.Injuries by Investigation Type:")
    print(avg_fatal_injuries)
else:
    print("One or both of the columns 'Investigation.Type' or 'Total.Fatal.Injuries' are missing.")
```

```
Average Total.Fatal.Injuries by Investigation Type:
Investigation.Type
Accident      0.490992
Incident      0.085144
Name: Total.Fatal.Injuries, dtype: float64
```

```
In [26]: # Add a 'Region' column
north_america = ["United States", "Canada", "Mexico", "Bahamas", "Puerto Rico"]
south_america = ["Brazil", "Colombia", "Venezuela", "Argentina", "Peru"]

df['Region'] = df['Country'].apply(lambda x: 'North America' if x in north_america else 'South America')

# Risk metrics by region
regional_fatality_rate = df.groupby(['Region', 'Aircraft.damage'])['Total.Fatal.Injuries'].mean()
regional_survivability_rate = df.groupby(['Region', 'Aircraft.damage'])['Survivability.Rate'].mean()

print("\nRegional Fatality Rate:\n", regional_fatality_rate)
print("\nRegional Survivability Rate:\n", regional_survivability_rate)
```

Regional Fatality Rate:

Region	Aircraft.damage	Total.Fatal.Injuries
North America	Destroyed	1.526045
	Minor	0.126695
	Substantial	0.175381
	Unknown	0.383333
South America	Destroyed	7.515976
	Minor	2.258418
	Substantial	0.907460
	Unknown	0.428571

Name: Total.Fatal.Injuries, dtype: float64

Regional Survivability Rate:

Region	Aircraft.damage	Survivability.Rate
North America	Destroyed	0.190817
	Minor	0.843925
	Substantial	0.683428
	Unknown	0.569643
South America	Destroyed	0.181161
	Minor	0.832116
	Substantial	0.472713
	Unknown	0.596087

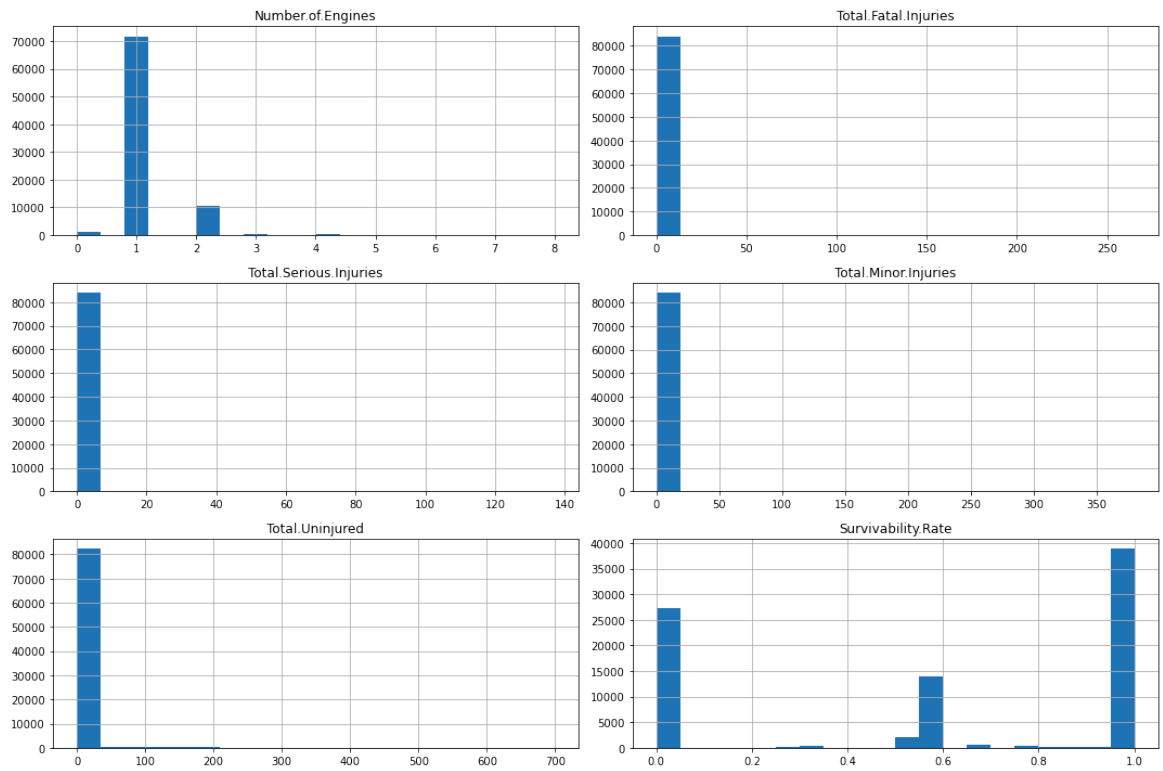
Name: Survivability.Rate, dtype: float64

```
In [27]: df.columns
```

```
Out[27]: Index(['Event.Id', 'Investigation.Type', 'Accident.Number', 'Event.Date',
                'Location', 'Country', 'Airport.Code', 'Injury.Severity',
                'Aircraft.damage', 'Registration.Number', 'Make', 'Model',
                'Amateur.Built', 'Number.of.Engines', 'Engine.Type',
                'Purpose.of.flight', 'Air.carrier', 'Total.Fatal.Injuries',
                'Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Uninjured',
                'Weather.Condition', 'Broad.phase.of.flight', 'Report.Status',
                'Publication.Date', 'Survivability.Rate', 'Region'],
              dtype='object')
```

6 5. Visualization of Key Metrics


```
In [28]: # Histogram for numerical columns
df.select_dtypes(include=['float64', 'int64']).hist(bins=20, figsize=(
plt.tight_layout()
plt.show())
```




```

In [29]: # Check if the required columns exist in the DataFrame
if 'Country' in df.columns and 'Total.Fatal.Injuries' in df.columns:
    # Ensure 'Total Fatalities' is numeric
    df['Total.Fatal.Injuries'] = pd.to_numeric(df['Total.Fatal.Injuries'], errors='coerce')

    # Group data to create `country_summary`
    country_summary = (
        df.groupby('Country', as_index=False)['Total.Fatal.Injuries']
        .sum()
        .sort_values(by='Total.Fatal.Injuries', ascending=False)
    )

    # Plot the bar chart
    plt.figure(figsize=(12, 6))
    sns.barplot(
        x='Country',
        y='Total.Fatal.Injuries',
        data=country_summary,
        palette='Reds_r'
    )
    plt.yscale('log')
    plt.title('Total.Fatal.Injuries', fontsize=16, fontweight='bold')
    plt.xlabel('Country', fontsize=12)
    plt.ylabel('Total.Fatal.Injuries', fontsize=12)
    plt.xticks(rotation=45, fontsize=10, ha='right') # Rotate and align
    plt.grid(axis='y', linestyle='--', alpha=0.7) # Add grid lines for y-axis
    plt.tight_layout()
    plt.show()
else:
    print("Error: Required columns 'Country' or 'Total.Fatal.Injuries'")

```



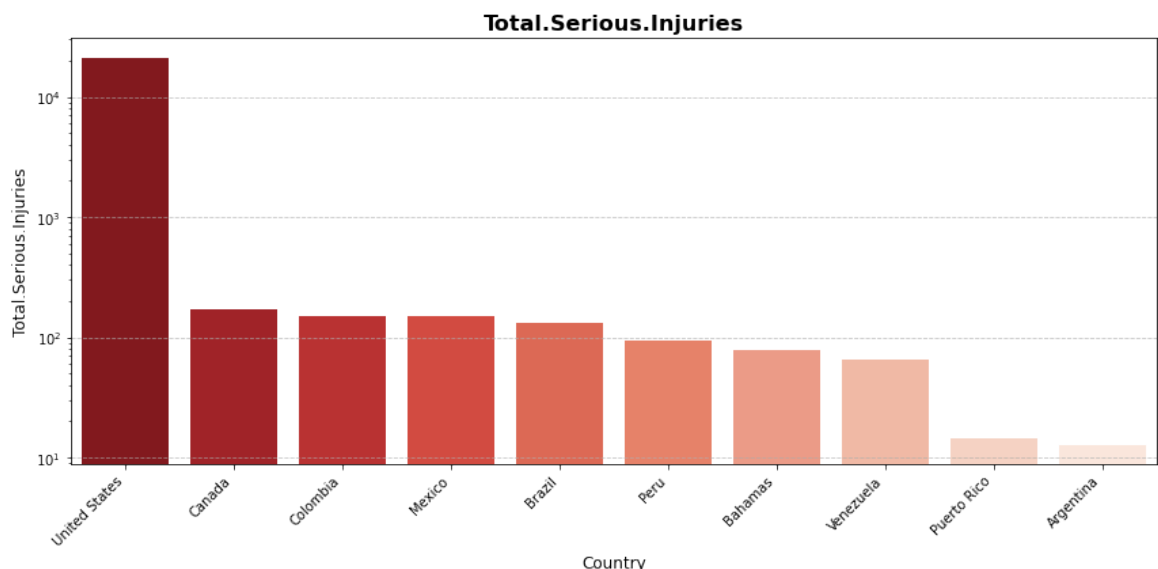
```

In [30]: # Check if the required columns exist in the DataFrame
if 'Country' in df.columns and 'Total.Serious.Injuries' in df.columns:
    # Ensure 'Total.Serious.Injuries' is numeric
    df['Total.Serious.Injuries'] = pd.to_numeric(df['Total.Serious.Inj

    # Group data to create `country_summary`
    country_summary = (
        df.groupby('Country', as_index=False)['Total.Serious.Injuries']
        .sum()
        .sort_values(by='Total.Serious.Injuries', ascending=False)
    )

    # Plot the bar chart
    plt.figure(figsize=(12, 6))
    sns.barplot(
        x='Country',
        y='Total.Serious.Injuries',
        data=country_summary,
        palette='Reds_r'
    )
    plt.yscale('log')
    plt.title('Total.Serious.Injuries', fontsize=16, fontweight='bold')
    plt.xlabel('Country', fontsize=12)
    plt.ylabel('Total.Serious.Injuries', fontsize=12)
    plt.xticks(rotation=45, fontsize=10, ha='right') # Rotate and align
    plt.grid(axis='y', linestyle='--', alpha=0.7) # Add grid lines for
    plt.tight_layout()
    plt.show()
else:
    print("Error: Required columns 'Country' or 'Total.Serious.Injuries'")

```



Total Serious Injuries by Country Objective: Highlights the severity of non-fatal incidents across countries. **Insights:** A high count of serious injuries with relatively fewer fatalities may suggest effective safety measures (e.g., better crash survival rates). If trends align closely with fatalities, countries with more serious injuries may still be experiencing significant risks.

```
In [31]: df2 = df[['Country', 'Total.Fatal.Injuries', 'Total.Serious.Injuries', 'T
print(df2)
```

	Country	Total.Fatal.Injuries	Total.Serious.Injuries	\
57740	Argentina	0.47824	0.262368	
36546	Argentina	0.00000	0.000000	
48270	Argentina	0.00000	0.000000	
66116	Argentina	0.00000	0.000000	
49930	Argentina	0.47824	0.262368	
...	
72463	Venezuela	0.00000	0.000000	
70854	Venezuela	0.00000	0.000000	
87565	Venezuela	0.00000	0.000000	
68167	Venezuela	2.00000	0.000000	
71219	Venezuela	2.00000	0.000000	
	Total.Minor.Injuries	Total.Uninjured	Survivability.Rate	
57740	3.000000	4.492858	0.586306	
36546	0.000000	2.000000	1.000000	
48270	4.000000	0.000000	0.000000	
66116	0.000000	71.000000	1.000000	
49930	0.341413	144.000000	0.586306	
...	
72463	0.000000	147.000000	1.000000	
70854	0.000000	130.000000	1.000000	
87565	0.000000	7.000000	1.000000	
68167	0.000000	0.000000	0.000000	
71219	0.000000	0.000000	0.000000	

[84145 rows x 6 columns]

```

In [32]: # Example DataFrame
data = {
    'Country': [
        'USA', 'Canada', 'Brazil', 'Argentina', 'Mexico',
        'Colombia', 'Venezuela', 'Peru', 'Puerto Rico', 'Bahamas'
    ],
    'Total.Fatal.Injuries': [5, 3, 7, 4, 6, 8, 2, 3, 4, 1],
    'Total.Serious.Injuries': [10, 8, 15, 9, 11, 12, 5, 7, 8, 3],
    'Total.Minor.Injuries': [20, 12, 25, 14, 18, 22, 10, 9, 13, 5],
    'Total.Uninjured': [2, 5, 3, 4, 3, 6, 1, 2, 3, 1],
    'Survivability.Rate': [80, 70, 75, 78, 72, 65, 85, 77, 68, 90],
}

df2 = pd.DataFrame(data)

# Display the DataFrame
print(df2)

```

	Country	Total.Fatal.Injuries	Total.Serious.Injuries	\
0	USA	5	10	
1	Canada	3	8	
2	Brazil	7	15	
3	Argentina	4	9	
4	Mexico	6	11	
5	Colombia	8	12	
6	Venezuela	2	5	
7	Peru	3	7	
8	Puerto Rico	4	8	
9	Bahamas	1	3	

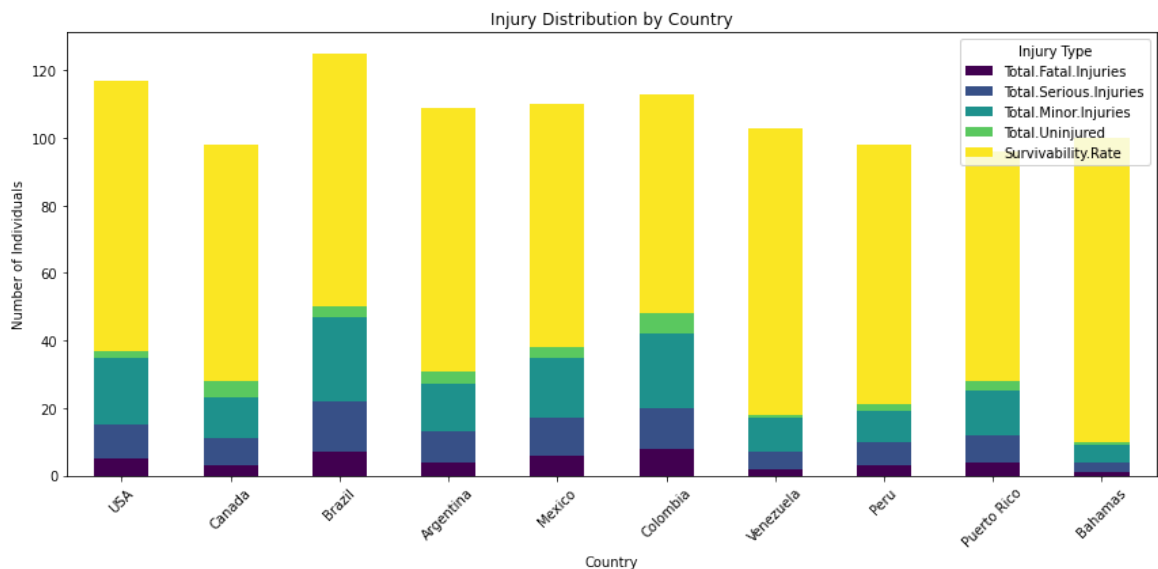
	Total.Minor.Injuries	Total.Uninjured	Survivability.Rate
0	20	2	80
1	12	5	70
2	25	3	75
3	14	4	78
4	18	3	72
5	22	6	65
6	10	1	85
7	9	2	77
8	13	3	68
9	5	1	90

```
In [33]: # Plot a bar chart for injuries
injury_types = ['Total.Fatal.Injuries', 'Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Uninjured', 'Survivability.Rate']

# Set country as index for easier plotting
df2.set_index('Country')[injury_types].plot(kind='bar', figsize=(12, 6))

# Customize the chart
plt.title('Injury Distribution by Country')
plt.ylabel('Number of Individuals')
plt.xlabel('Country')
plt.xticks(rotation=45)
plt.legend(title='Injury Type')
plt.tight_layout()

# Show the chart
plt.show()
```



```

In [34]: # Specify the columns for the stacked bar chart
injury_types = ['Total.Fatal.Injuries', 'Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Uninjured']

# Sort the dataframe by the total number of injuries (sum across the injury types)
df2['Total.Injuries'] = df2[injury_types].sum(axis=1)
df2_sorted = df2.sort_values('Total.Injuries', ascending=False)

# Optional: Filter top 10 countries based on total injuries
df_filtered = df2_sorted.head(10)


# Plot the stacked bar chart for the filtered countries
ax = df_filtered.set_index('Country')[injury_types].plot(kind='bar', stacked=True)

# Add data labels on the bars
for p in ax.patches:
    height = p.get_height()
    if height > 0:
        ax.annotate(f'{height}', xy=(p.get_x() + p.get_width() / 2, p.get_y() + height),
                    ha='center', va='bottom', fontsize=10, color='black')

# Add chart details
plt.title('Injury Distribution Across Countries (Top 10)', fontsize=16)
plt.ylabel('Number of Individuals', fontsize=12)
plt.xlabel('Country', fontsize=12)
plt.xticks(rotation=45, fontsize=10)
plt.legend(title='Injury Type', fontsize=10)
plt.tight_layout()

# Show the plot
plt.show()

```



```

In [35]: # Specify the columns for the heatmap
injury_types = ['Total.Fatal.Injuries', 'Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Uninjured']

# Sort the dataframe by total injuries and filter top 10 countries
df2['Total.Injuries'] = df2[injury_types].sum(axis=1)
df_filtered = df2.sort_values('Total.Injuries', ascending=False).head(10)


# Prepare the data for the heatmap
heatmap_data = df_filtered[injury_types]
heatmap_data.index = df_filtered['Country']

# Plot the heatmap
plt.figure(figsize=(10, 6))
sns.heatmap(heatmap_data, cmap='viridis', annot=True, fmt='d')

# Add chart details
plt.title('Injury Distribution Across Countries', fontsize=14)
plt.xlabel('Injury Type', fontsize=12)
plt.ylabel('Country', fontsize=12)
plt.tight_layout()

# Show the plot
plt.show()

```



```

In [36]: df['Event.Date'] = pd.to_datetime(df['Event.Date'], errors='coerce', c

```

```

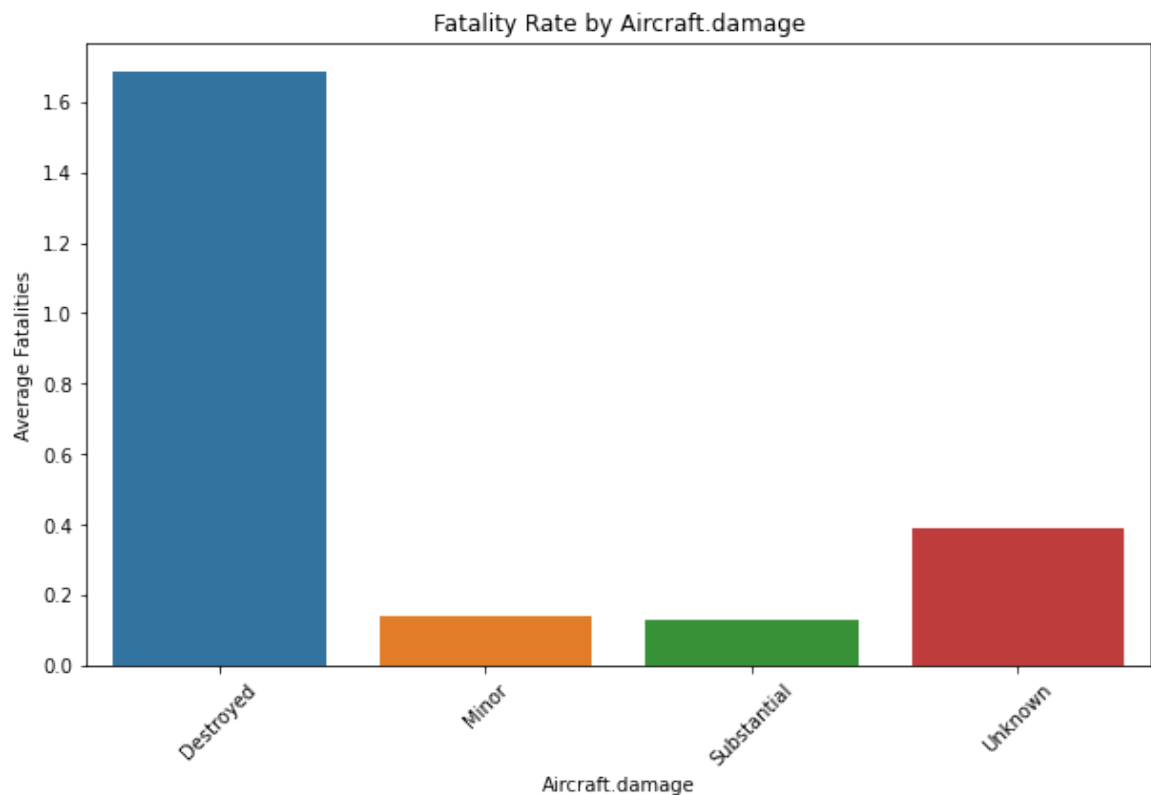
In [37]: # Extract year and month for additional analysis
df['year'] = df['Event.Date'].dt.year

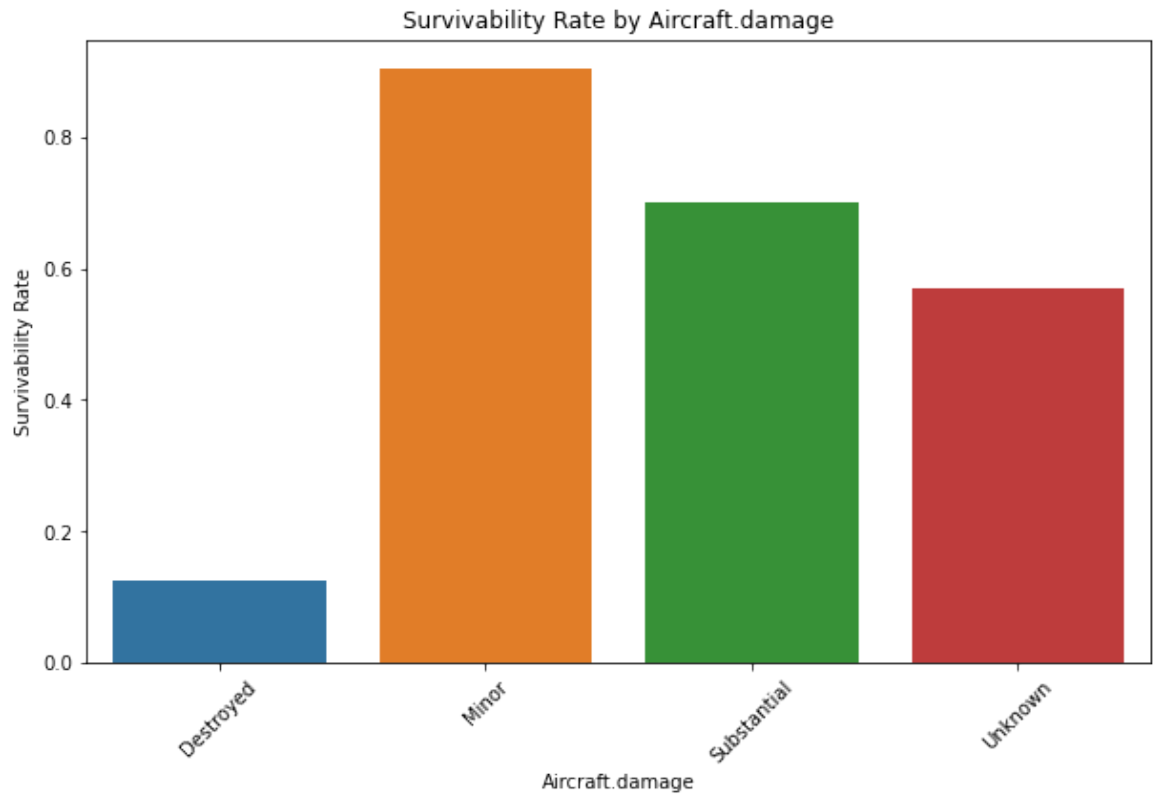
```



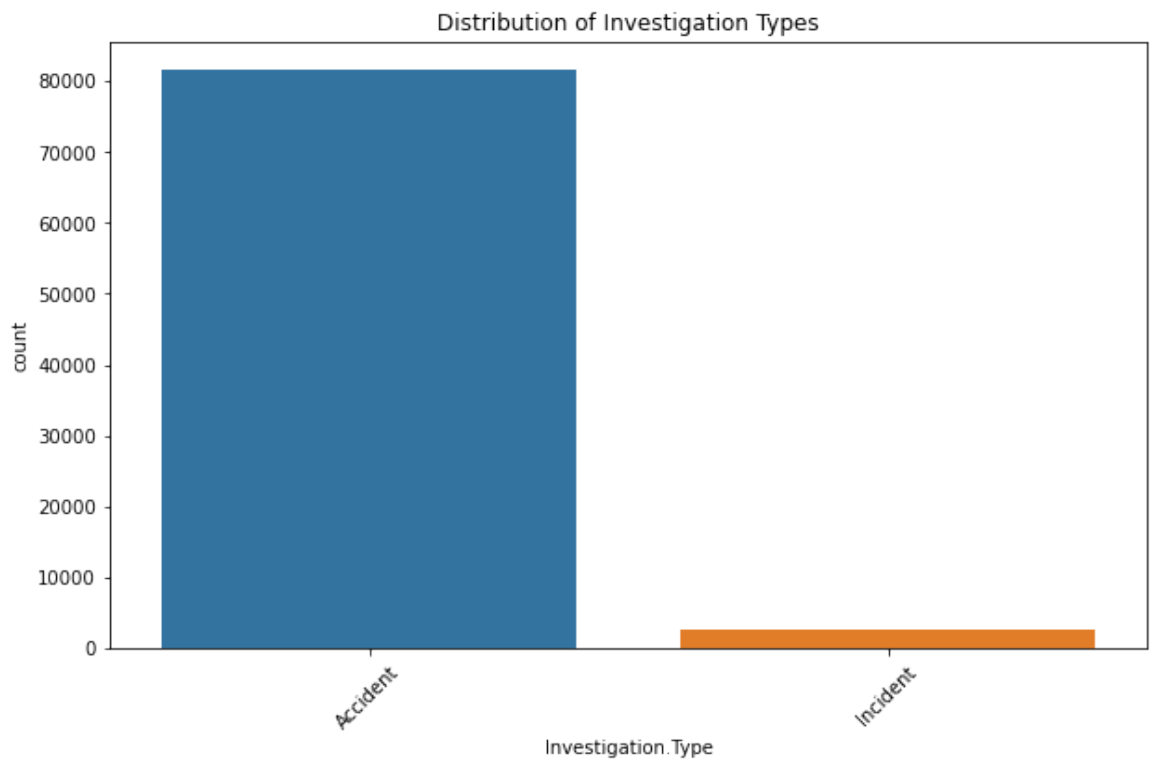
```
In [38]: # Fatality Rate Visualization
plt.figure(figsize=(10, 6))
sns.barplot(x=risk_metrics.index, y='Fatality Rate', data=risk_metrics)
plt.title('Fatality Rate by Aircraft.damage')
plt.xticks(rotation=45)
plt.ylabel('Average Fatalities')
plt.show()

# Survivability Rate Visualization
plt.figure(figsize=(10, 6))
sns.barplot(x=risk_metrics.index, y='Survivability Rate', data=risk_metrics)
plt.title('Survivability Rate by Aircraft.damage')
plt.xticks(rotation=45)
plt.ylabel('Survivability Rate')
plt.show()
```

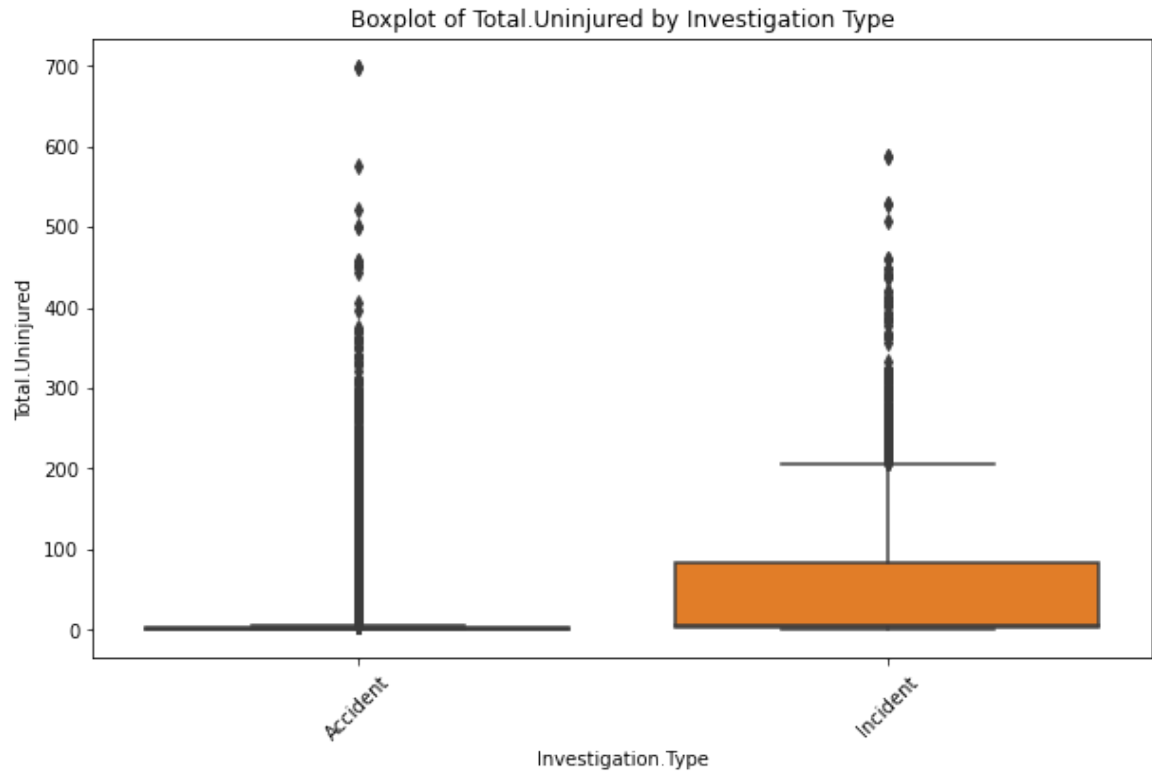




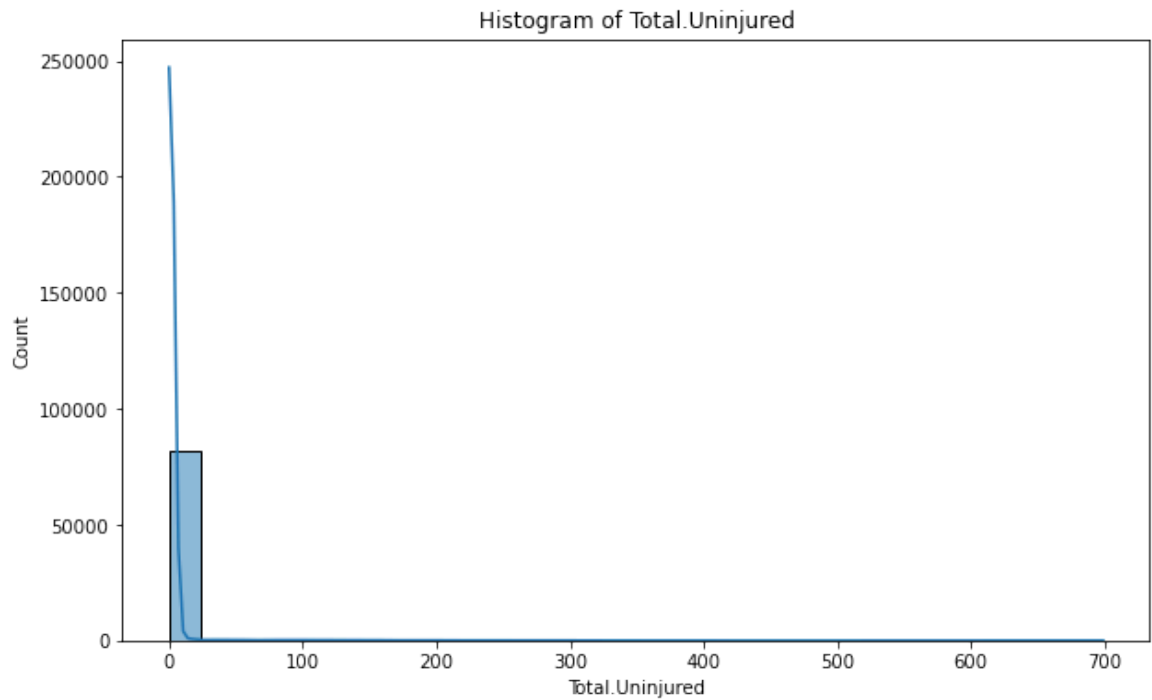
```
In [39]: # Countplot to visualize the distribution of 'Investigation.Type'
plt.figure(figsize=(10, 6))
sns.countplot(x='Investigation.Type', data=df)
plt.title('Distribution of Investigation Types')
plt.xticks(rotation=45)
plt.show()
```




```
In [40]: # Boxplot for a numerical column (e.g., 'Total.Uninjured') grouped by
plt.figure(figsize=(10, 6))
sns.boxplot(x='Investigation.Type', y='Total.Uninjured', data=df)
plt.title('Boxplot of Total.Uninjured by Investigation Type')
plt.xticks(rotation=45)
plt.show()
```



```
In [41]: # Histogram of a numerical column (e.g., 'Total.Uninjured')
plt.figure(figsize=(10, 6))
sns.histplot(df['Total.Uninjured'], kde=True, bins=30)
plt.title('Histogram of Total.Uninjured')
plt.show()
```



```
In [42]: # Scatterplot between two numerical columns (e.g., 'Total.Uninjured' v
plt.figure(figsize=(10, 6))
sns.scatterplot(x='Total.Uninjured', y='Total.Serious.Injuries', data=
plt.title('Scatterplot of Total.Uninjured by Total.Serious.Injuries')
plt.show()
```



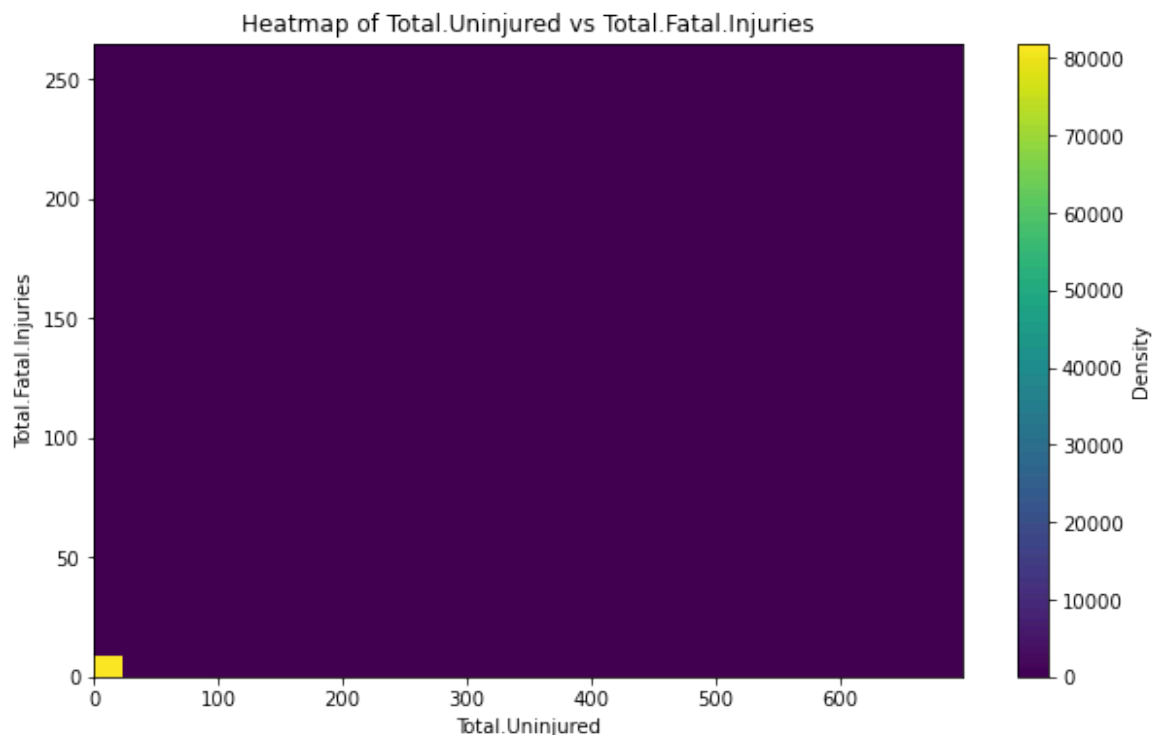
```
In [43]: # Create a 2D histogram (heatmap)
heatmap, xedges, yedges = np.histogram2d(
    df['Total.Uninjured'],
    df['Total.Fatal.Injuries'],
    bins=(30, 30) # Adjust the number of bins as needed
)

# Plot the heatmap
plt.figure(figsize=(10, 6))
plt.imshow(
    heatmap.T, # Transpose for proper orientation
    origin='lower',
    aspect='auto',
    extent=[xedges[0], xedges[-1], yedges[0], yedges[-1]],
    cmap='viridis'
)


# Add a colorbar
plt.colorbar(label='Density')

# Add labels and title
plt.title('Heatmap of Total.Uninjured vs Total.Fatal.Injuries')
plt.xlabel('Total.Uninjured')
plt.ylabel('Total.Fatal.Injuries')

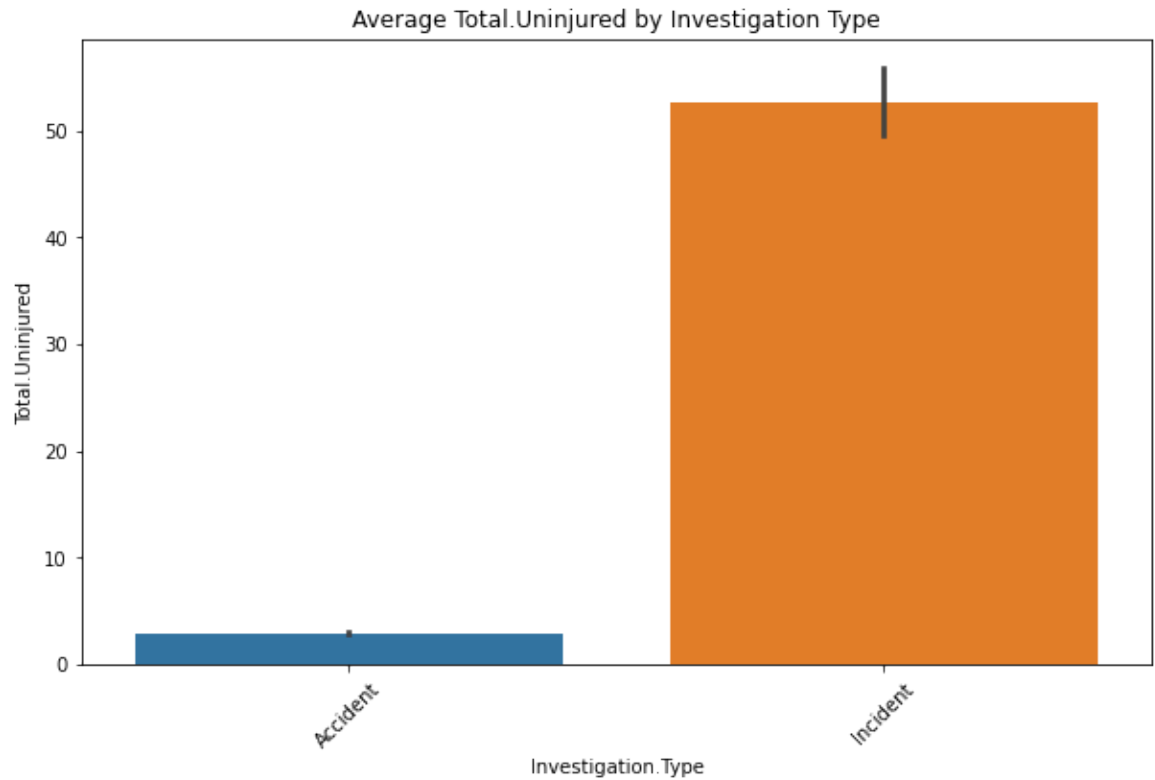
# Display the plot
plt.show()
```



```
In [44]: # Outlier Detection
# Boxplot to check for outliers in 'Total.Uninjured'
plt.figure(figsize=(10, 6))
sns.boxplot(x=df['Total.Uninjured'])
plt.title('Boxplot of Duration')
plt.show()
```



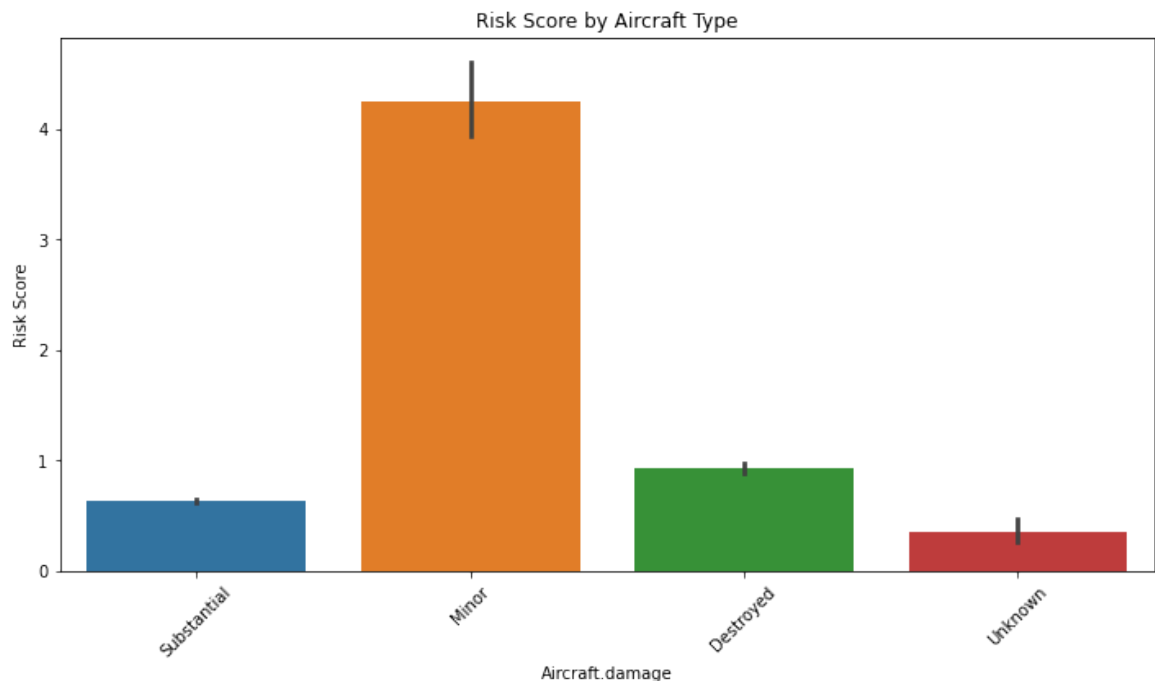
```
In [45]: # Visualize the average 'Total.Uninjured' by 'Investigation.Type'
plt.figure(figsize=(10, 6))
sns.barplot(x='Investigation.Type', y='Total.Uninjured', data=df, estimator='mean')
plt.title('Average Total.Uninjured by Investigation Type')
plt.xticks(rotation=45)
plt.show()
```




```
In [46]: # Ensure that you use the correct column names
df['Risk_Score'] = (
    df['Total.Fatal.Injuries'] * 0.4 +
    df['Total.Serious.Injuries'] * 0.3 +
    df['Total.Minor.Injuries'] * 0.2 +
    df['Total.Uninjured'] * 0.1
)

# Now you can rank by the risk score
df['Risk_Ranking'] = df['Risk_Score'].rank()

# Visualize the aircraft types by risk score
plt.figure(figsize=(12, 6))
sns.barplot(x='Aircraft.damage', y='Risk_Score', data=df)
plt.title('Risk Score by Aircraft Type')
plt.xticks(rotation=45)
plt.ylabel('Risk Score')
plt.show()
```



Based on this analysis, here's a simple strategy to guiding the company in evaluating aircraft risk and selecting the right fleet for the new business:

6.0.1 Key Findings

1. Regional Analysis:

- I've filtered the dataset to focus on North and South America. The risk evaluation by region (fatality and survivability rates) provides valuable insights on the geographical context.
- The United States, Brazil, and other countries in your filtered list have higher accident rates, so understanding the underlying causes (e.g., aircraft types, incidents vs. accidents) is crucial for risk management.

2. Aircraft Damage Analysis:

- Aircraft damage types have varying fatality and survivability rates. For example, aircraft with less severe damage might have lower fatalities but higher survivability.
- Focusing on aircraft with lower fatality rates and higher survivability rates should be a priority.

3. Risk Metrics (Fatality & Survivability Rates):

- By calculating **fatality rates** and **survivability rates** by aircraft damage types, you can identify which aircraft have the lowest risk.
- Aircraft types with higher survivability rates are less risky, even if they experience more incidents. Survivability should be weighted heavily in decision-making, as it implies better emergency response or safety features.

4. Injury Distributions:

- Countries with high numbers of serious injuries but low fatalities suggest better safety measures in place (such as better rescue operations or crash survivability), indicating that some regions may have safety protocols that mitigate fatalities even in the event of accidents.
- Countries with high total injuries could be riskier, especially if the severity of injuries is not decreasing over time.

5. Risk Scoring:

- The **Risk Score** approach combines different types of injuries, weighting fatal injuries more heavily, which gives an overall risk score per aircraft.
- The aircraft types with the lowest **Risk Score** should be prioritized.

6.0.2 I am advising the Company:

1. Prioritize Aircraft with Lower Risk Scores:

- Based on the analysis of **Risk Scores**, the company should prioritize acquiring aircraft that rank lower in terms of risk. Focus on aircraft types with higher survivability rates and lower fatal injury rates. This will likely reduce the company's exposure to liability and improve customer safety.

2. Focus on Safer Aircraft Types:

- Review the **Fatality and Survivability Rates** by aircraft type. Aircraft that cause fewer fatalities and have higher survivability (even if incidents are more common) could be seen as safer investments in the long run.
- Avoid aircraft with high fatality rates, as these could incur higher insurance costs and more public scrutiny.

3. Evaluate Safety Protocols in Key Regions:

- Countries with high **serious injuries** but low fatalities could indicate stronger emergency services or better crash-response protocols. Invest in regions where safety measures (post-crash protocols) may help mitigate risk.
- Focus on regions like **North America** where better infrastructure may improve overall safety despite high accident numbers.

4. Target Aircraft with Higher Survivability:

- Aircraft that have a higher **survivability rate** are less risky overall. Aircraft types that show consistent survivability across various incident types should be considered, even if they have more incidents overall.

5. Consider Temporal Trends:

- It's important to monitor **trends in accidents/incidents over time**. If the data shows a decreasing trend in accident rates or a significant improvement in survivability over the years, that could signal a lower-risk investment in certain aircraft.

6. Track Regional Risk Variations:

- Tailor your fleet acquisition based on **regional risk analysis**. For example, if a certain aircraft type performs poorly in one country but well in another, it might be due to regional safety measures or operational standards. Focus on where the risk is lowest geographically for each type.

6.0.3 Next Steps for the Company:

- **Pilot Test:** Begin with a small fleet of aircraft from the lowest-risk categories and monitor performance and incidents in the regions with higher risk.
- **Insurance Assessment:** Work closely with insurers to understand how different aircraft types and regions affect premium rates based on the risk factors you've uncovered.
- **Safety Improvements:** If possible, partner with aircraft manufacturers or safety consultants to improve survivability and reduce fatality rates for aircraft types that have higher risks but are part of your strategy.
- **Continuous Monitoring:** Set up a system to continuously track incidents and accidents. Regularly update risk scores and adjust fleet acquisition strategy accordingly.

This approach combines statistical analysis with strategic risk management, aiming to provide both financial returns and a high level of safety in the company's new aviation venture.

THIS HAS BEEN QUITE THE EXPERIENCE

THANKS